

# exp-ml-1

October 14, 2024

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor

from sklearn.metrics import r2_score, mean_squared_error
```

```
[7]: data = pd.read_csv("uber.csv")
data
```

```
[7]:
```

	Unnamed: 0	key	fare_amount	\
0	24238194	2015-05-07 19:52:06.0000003	7.5	
1	27835199	2009-07-17 20:04:56.0000002	7.7	
2	44984355	2009-08-24 21:45:00.00000061	12.9	
3	25894730	2009-06-26 08:22:21.0000001	5.3	
4	17610152	2014-08-28 17:47:00.000000188	16.0	
...	...	...	...	
199995	42598914	2012-10-28 10:49:00.00000053	3.0	
199996	16382965	2014-03-14 01:09:00.0000008	7.5	
199997	27804658	2009-06-29 00:42:00.00000078	30.9	
199998	20259894	2015-05-20 14:56:25.0000004	14.5	
199999	11951496	2010-05-15 04:08:00.00000076	14.1	

	pickup_datetime	pickup_longitude	pickup_latitude	\
0	2015-05-07 19:52:06 UTC	-73.999817	40.738354	
1	2009-07-17 20:04:56 UTC	-73.994355	40.728225	
2	2009-08-24 21:45:00 UTC	-74.005043	40.740770	
3	2009-06-26 08:22:21 UTC	-73.976124	40.790844	
4	2014-08-28 17:47:00 UTC	-73.925023	40.744085	
...	...	...	...	
199995	2012-10-28 10:49:00 UTC	-73.987042	40.739367	
199996	2014-03-14 01:09:00 UTC	-73.984722	40.736837	
199997	2009-06-29 00:42:00 UTC	-73.986017	40.756487	
199998	2015-05-20 14:56:25 UTC	-73.997124	40.725452	

199999 2010-05-15 04:08:00 UTC -73.984395 40.720077

	dropoff_longitude	dropoff_latitude	passenger_count
0	-73.999512	40.723217	1
1	-73.994710	40.750325	1
2	-73.962565	40.772647	1
3	-73.965316	40.803349	3
4	-73.973082	40.761247	5
...	...	...	...
199995	-73.986525	40.740297	1
199996	-74.006672	40.739620	1
199997	-73.858957	40.692588	2
199998	-73.983215	40.695415	1
199999	-73.985508	40.768793	1

[200000 rows x 9 columns]

```
[8]: # 1. Pre-process the dataset

# Remove unnecessary column
data["pickup_datetime"] = pd.to_datetime(data["pickup_datetime"])

missing_values = data.isnull().sum()
print("Missing values in the dataset:")
print(missing_values)

# Handle missing values
# We can choose to drop rows with missing values or fill them with appropriate
↪ values.

data.dropna(inplace=True)

# To fill missing values with the mean value of the column:
# data.fillna(data.mean(), inplace=True)

# Ensure there are no more missing values
missing_values = data.isnull().sum()
print("Missing values after handling:")
print(missing_values)

# 2. Identify outliers
# visualization to detect outliers.
sns.boxplot(x=data["fare_amount"])
plt.show()
```


Missing values in the dataset:

Unnamed: 0 0

```

key                0
fare_amount        0
pickup_datetime    0
pickup_longitude    0
pickup_latitude     0
dropoff_longitude   1
dropoff_latitude    1
passenger_count     0
dtype: int64
Missing values after handling:
Unnamed: 0         0
key                0
fare_amount        0
pickup_datetime    0
pickup_longitude    0
pickup_latitude     0
dropoff_longitude   0
dropoff_latitude    0
passenger_count     0
dtype: int64

```




```
[9]: # Calculate the IQR for the 'fare_amount' column
Q1 = data["fare_amount"].quantile(0.25)
Q3 = data["fare_amount"].quantile(0.75)
IQR = Q3 - Q1

# Define a threshold (e.g., 1.5 times the IQR) to identify outliers
threshold = 1.5
lower_bound = Q1 - threshold * IQR
upper_bound = Q3 + threshold * IQR

# Remove outliers
data_no_outliers = data[(data["fare_amount"] >= lower_bound) &
    ↪ (data["fare_amount"] <= upper_bound)]

# Visualize the 'fare_amount' distribution without outliers
sns.boxplot(x=data_no_outliers["fare_amount"])
plt.show()
```




```
[10]: data.plot(kind="box",subplots=True, layout=(7, 2), figsize=(15, 20))
```

```

[10]: Unnamed: 0      Axes(0.125,0.786098;0.352273x0.0939024)
      fare_amount     Axes(0.547727,0.786098;0.352273x0.0939024)
      pickup_longitude Axes(0.125,0.673415;0.352273x0.0939024)
      pickup_latitude  Axes(0.547727,0.673415;0.352273x0.0939024)
      dropoff_longitude Axes(0.125,0.560732;0.352273x0.0939024)
      dropoff_latitude  Axes(0.547727,0.560732;0.352273x0.0939024)
      passenger_count   Axes(0.125,0.448049;0.352273x0.0939024)
      dtype: object

```



```

[12]: # 4. Implement linear regression and random forest regression models
      # Split the data into features and target variable
      X = data[['pickup_longitude', 'pickup_latitude', 'dropoff_longitude', '
      ↪ 'dropoff_latitude', 'passenger_count']]
      y = data['fare_amount'] #Target
      y

```

```

[12]: 0      7.5
      1      7.7
      2     12.9
      3      5.3

```

```

4          16.0
...
199995     3.0
199996     7.5
199997    30.9
199998    14.5
199999    14.1
Name: fare_amount, Length: 199999, dtype: float64

```

```

[13]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)

```

```

[14]: # Create and train the linear regression model
lr_model = LinearRegression()
lr_model.fit(X_train, y_train)

```

```

[14]: LinearRegression()

```

```

[ ]: # Create and train the random forest regression model
rf_model = RandomForestRegressor(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)

```

```

[ ]: # 5. Evaluate the models
# Predict the values
y_pred_lr = lr_model.predict(X_test)
y_pred_lr
print("Linear Model:", y_pred_lr)
y_pred_rf = rf_model.predict(X_test)
print("Random Forest Model:", y_pred_rf)

```

```

[ ]: # Calculate R-squared (R2) and Root Mean Squared Error (RMSE) for both models
r2_lr = r2_score(y_test, y_pred_lr)
rmse_lr = np.sqrt(mean_squared_error(y_test, y_pred_lr))

```

```

[ ]: # Compare the scores
print("Linear Regression - R2:", r2_lr)
print("Linear Regression - RMSE:", rmse_lr)

```

```

[ ]: # Compare the scores
print("Linear Regression - R2:", r2_lr)
print("Linear Regression - RMSE:", rmse_lr)

```

```

[ ]:

```

# exp-ml-2

October 14, 2024

```
[1]: # Import necessary libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, classification_report
```

```
[2]: # Load the dataset
data = pd.read_csv("emails.csv") # Replace with the actual path to the dataset
data
```

```
[2]:
```

	Email No.	the	to	ect	and	for	of	a	you	hou	...	connevey	\
0	Email 1	0	0	1	0	0	0	2	0	0	...	0	
1	Email 2	8	13	24	6	6	2	102	1	27	...	0	
2	Email 3	0	0	1	0	0	0	8	0	0	...	0	
3	Email 4	0	5	22	0	5	1	51	2	10	...	0	
4	Email 5	7	6	17	1	5	2	57	0	9	...	0	
...	...	...	...	...	...	...	...	...	...	...	...	...	
5167	Email 5168	2	2	2	3	0	0	32	0	0	...	0	
5168	Email 5169	35	27	11	2	6	5	151	4	3	...	0	
5169	Email 5170	0	0	1	1	0	0	11	0	0	...	0	
5170	Email 5171	2	7	1	0	2	1	28	2	0	...	0	
5171	Email 5172	22	24	5	1	6	5	148	8	2	...	0	

	jay	valued	lay	infrastructure	military	allowing	ff	dry	\
0	0	0	0		0	0	0	0	
1	0	0	0		0	0	0	1	0
2	0	0	0		0	0	0	0	0
3	0	0	0		0	0	0	0	0
4	0	0	0		0	0	0	1	0
...	...	...	...	...	...	...	...	...	...
5167	0	0	0		0	0	0	0	0
5168	0	0	0		0	0	0	1	0
5169	0	0	0		0	0	0	0	0
5170	0	0	0		0	0	0	1	0
5171	0	0	0		0	0	0	0	0

	Prediction
0	0
1	0
2	0
3	0
4	0
...	...
5167	0
5168	0
5169	1
5170	1
5171	0

[5172 rows x 3002 columns]

```
[3]: # 1. Data Preprocessing - Handle missing values if necessary
data.drop(['Email No.'],axis=1, inplace=True)
# 2. Feature Selection/Engineering - Select relevant features
```

```
[4]: # 3. Split the data into training and testing sets
X = data.drop("Prediction", axis=1) # Features
y = data["Prediction"] # Target variable
print("Features: ",X)
print("Target: ",y)
```

Features:		the	to	ect	and	for	of	a	you	hou	in	...
enhancements \												
0	0	0	1	0	0	0	2	0	0	0	...	0
1	8	13	24	6	6	2	102	1	27	18	...	0
2	0	0	1	0	0	0	8	0	0	4	...	0
3	0	5	22	0	5	1	51	2	10	1	...	0
4	7	6	17	1	5	2	57	0	9	3	...	0
...	...	...	...	...	...	...	...	...	...	...	...	...
5167	2	2	2	3	0	0	32	0	0	5	...	0
5168	35	27	11	2	6	5	151	4	3	23	...	0
5169	0	0	1	1	0	0	11	0	0	1	...	0
5170	2	7	1	0	2	1	28	2	0	8	...	0
5171	22	24	5	1	6	5	148	8	2	23	...	0

	connevey	jay	valued	lay	infrastructure	military	allowing	ff	dry
0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	1	0
2	0	0	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	1	0
...	...	...	...	...	...	...	...	...	...
5167	0	0	0	0	0	0	0	0	0



5168	0	0	0	0	0	0	0	0	1	0
5169	0	0	0	0	0	0	0	0	0	0
5170	0	0	0	0	0	0	0	0	1	0
5171	0	0	0	0	0	0	0	0	0	0

[5172 rows x 3000 columns]

Target: 0 0

1 0

2 0

3 0

4 0

..

5167 0

5168 0

5169 1

5170 1

5171 0

Name: Prediction, Length: 5172, dtype: int64

```
[5]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
↳ random_state=42)
```

```
[6]: # 4. Model Building
# K-Nearest Neighbors
knn_model = KNeighborsClassifier(n_neighbors=5)
knn_model.fit(X_train, y_train)

# Support Vector Machine
svm_model = SVC()
svm_model.fit(X_train, y_train)
```

```
[6]: SVC()
```

```
[7]: # 5. Model Evaluation
# K-Nearest Neighbors
knn_predictions = knn_model.predict(X_test)
knn_accuracy = accuracy_score(y_test, knn_predictions)
knn_report = classification_report(y_test, knn_predictions)
```

```
[8]: print(knn_predictions)
```

[0 0 1 ... 0 0 0]

```
[9]: # Print or visualize the evaluation results
print("K-Nearest Neighbors Accuracy:")
print(knn_accuracy)
print("K-Nearest Neighbors Classification Report:")
```

```
print(knn_report)
```

K-Nearest Neighbors Accuracy:

0.8608247422680413

K-Nearest Neighbors Classification Report:

	precision	recall	f1-score	support
0	0.93	0.87	0.90	1097
1	0.73	0.83	0.78	455
accuracy			0.86	1552
macro avg	0.83	0.85	0.84	1552
weighted avg	0.87	0.86	0.86	1552

```
[10]: # Support Vector Machine
svm_predictions = svm_model.predict(X_test)
svm_accuracy = accuracy_score(y_test, svm_predictions)
svm_report = classification_report(y_test, svm_predictions)
```

```
[11]: print(svm_predictions)
```

[0 0 1 ... 0 0 0]

```
[13]: print("Support Vector Machine Accuracy:")
print(svm_accuracy)
print("Support Vector Machine Classification Report:")
print(svm_report)
```

Support Vector Machine Accuracy:

0.803479381443299

Support Vector Machine Classification Report:

	precision	recall	f1-score	support
0	0.79	0.99	0.88	1097
1	0.92	0.36	0.52	455
accuracy			0.80	1552
macro avg	0.85	0.67	0.70	1552
weighted avg	0.83	0.80	0.77	1552

```
[ ]:
```

## exp-ml-3

October 14, 2024

```
[21]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, confusion_matrix
import tensorflow as tf
from tensorflow import keras
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[4]: # 1. Read the dataset
data = pd.read_csv("Churn_Modelling.csv") # Replace with the actual path to
↳ the dataset
data
```

```
[4]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	\
0	1	15634602	Hargrave	619	France	Female	42	
1	2	15647311	Hill	608	Spain	Female	41	
2	3	15619304	Onio	502	France	Female	42	
3	4	15701354	Boni	699	France	Female	39	
4	5	15737888	Mitchell	850	Spain	Female	43	
...	...	...	...	...	...	...	...	
9995	9996	15606229	Obijiaku	771	France	Male	39	
9996	9997	15569892	Johnstone	516	France	Male	35	
9997	9998	15584532	Liu	709	France	Female	36	
9998	9999	15682355	Sabbatini	772	Germany	Male	42	
9999	10000	15628319	Walker	792	France	Female	28	

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	
...	...	...	...	...	...	
9995	5	0.00	2	1	0	
9996	10	57369.61	1	1	1	
9997	7	0.00	1	0	1	

9998	3	75075.31	2	1	0
9999	4	130142.79	1	1	0

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0
...	...	...
9995	96270.64	0
9996	101699.77	0
9997	42085.58	1
9998	92888.52	1
9999	38190.78	0

[10000 rows x 14 columns]

```
[5]: # 2. Distinguish features and target
X = data.drop("Exited", axis=1) # Features
y = data["Exited"] # Target variable
```

```
[6]: X
```

```
[6]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	\
0	1	15634602	Hargrave	619	France	Female	42	
1	2	15647311	Hill	608	Spain	Female	41	
2	3	15619304	Onio	502	France	Female	42	
3	4	15701354	Boni	699	France	Female	39	
4	5	15737888	Mitchell	850	Spain	Female	43	
...	...	...	...	...	...	...	...	
9995	9996	15606229	Obijiaku	771	France	Male	39	
9996	9997	15569892	Johnstone	516	France	Male	35	
9997	9998	15584532	Liu	709	France	Female	36	
9998	9999	15682355	Sabbatini	772	Germany	Male	42	
9999	10000	15628319	Walker	792	France	Female	28	

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	
...	...	...	...	...	...	
9995	5	0.00	2	1	0	
9996	10	57369.61	1	1	1	
9997	7	0.00	1	0	1	

9998	3	75075.31	2	1	0
9999	4	130142.79	1	1	0

	EstimatedSalary
0	101348.88
1	112542.58
2	113931.57
3	93826.63
4	79084.10
...	...
9995	96270.64
9996	101699.77
9997	42085.58
9998	92888.52
9999	38190.78

[10000 rows x 13 columns]

```
[7]: # 2. Divide the dataset into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↪random_state=42)
```

```
[11]: data.head()
```

```
[11]:
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	\
0	619	France	Female	42	2	0.00		1
1	608	Spain	Female	41	1	83807.86		1
2	502	France	Female	42	8	159660.80		3
3	699	France	Female	39	1	0.00		2
4	850	Spain	Female	43	2	125510.82		1

	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1		101348.88	1
1	0		112542.58	0
2	1		113931.57	1
3	0		93826.63	0
4	1		79084.10	0

```
[12]: data.isna().any()
data.isna().sum()
```

```
[12]: CreditScore      0
Geography              0
Gender                 0
Age                   0
Tenure                 0
Balance                0
```

```

NumOfProducts      0
HasCrCard           0
IsActiveMember      0
EstimatedSalary     0
Exited              0
dtype: int64

```

```
[13]: print(data.shape)
      data.info()
```

```

(10000, 11)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 11 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   CreditScore           10000 non-null  int64
 1   Geography             10000 non-null  object
 2   Gender                10000 non-null  object
 3   Age                   10000 non-null  int64
 4   Tenure                10000 non-null  int64
 5   Balance               10000 non-null  float64
 6   NumOfProducts         10000 non-null  int64
 7   HasCrCard             10000 non-null  int64
 8   IsActiveMember        10000 non-null  int64
 9   EstimatedSalary        10000 non-null  float64
10   Exited                10000 non-null  int64
dtypes: float64(2), int64(7), object(2)
memory usage: 859.5+ KB

```

```
[14]: data.describe()
```

```
[14]:
```

	CreditScore	Age	Tenure	Balance	NumOfProducts \
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	650.528800	38.921800	5.012800	76485.889288	1.530200
std	96.653299	10.487806	2.892174	62397.405202	0.581654
min	350.000000	18.000000	0.000000	0.000000	1.000000
25%	584.000000	32.000000	3.000000	0.000000	1.000000
50%	652.000000	37.000000	5.000000	97198.540000	1.000000
75%	718.000000	44.000000	7.000000	127644.240000	2.000000
max	850.000000	92.000000	10.000000	250898.090000	4.000000

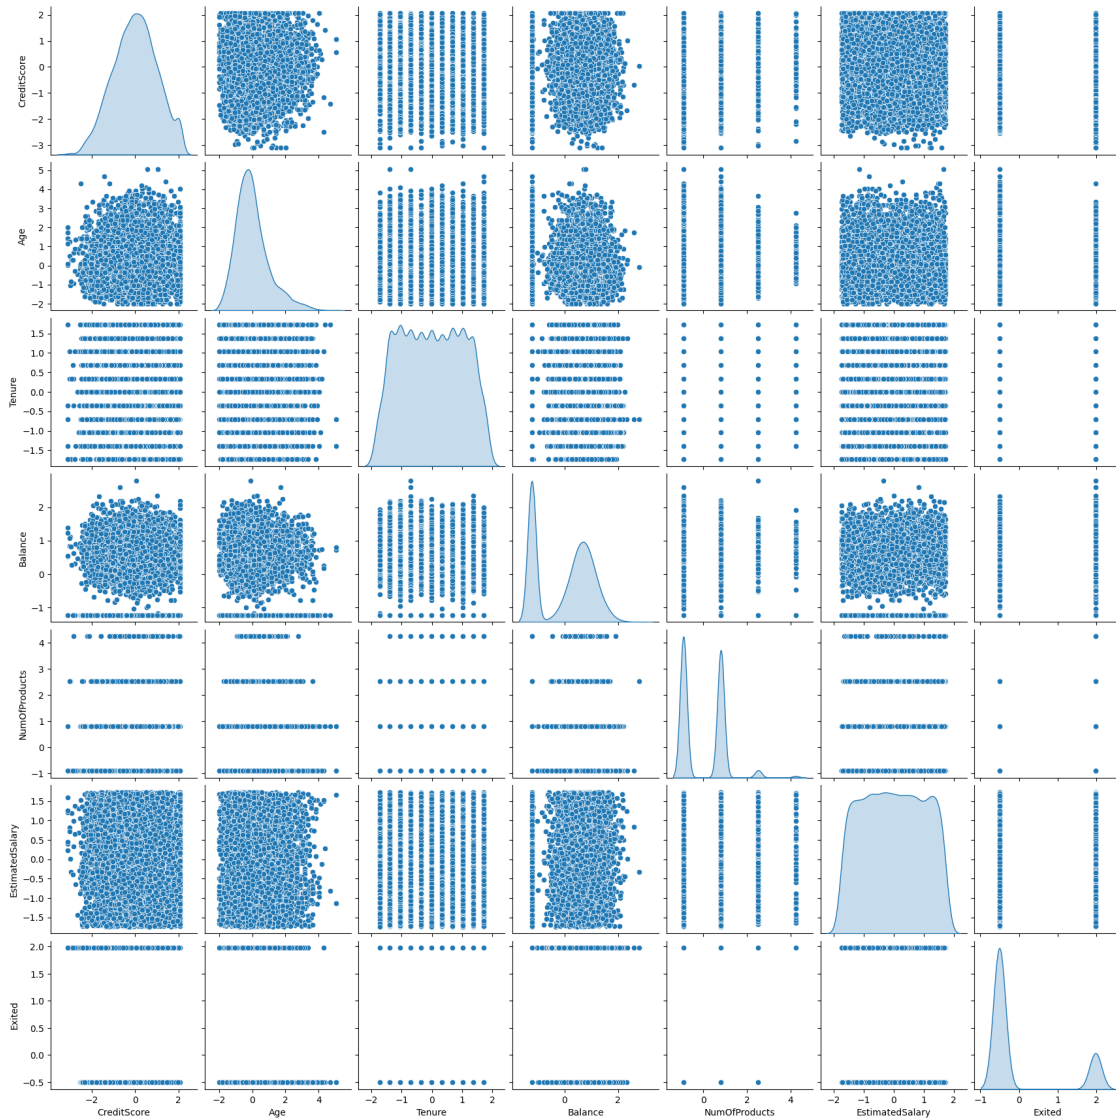
  

	HasCrCard	IsActiveMember	EstimatedSalary	Exited
count	10000.000000	10000.000000	10000.000000	10000.000000
mean	0.70550	0.515100	100090.239881	0.203700
std	0.45584	0.499797	57510.492818	0.402769
min	0.00000	0.000000	11.580000	0.000000

25%	0.00000	0.000000	51002.110000	0.000000
50%	1.00000	1.000000	100193.915000	0.000000
75%	1.00000	1.000000	149388.247500	0.000000
max	1.00000	1.000000	199992.480000	1.000000


```
[17]: ## Scale the data
scaler=StandardScaler()
## Extract only the Numerical Columns to perform Bivariate Analysis
subset=data.drop(['Geography','Gender','HasCrCard','IsActiveMember'],axis=1)
scaled=scaler.fit_transform(subset)
scaled_df=pd.DataFrame(scaled,columns=subset.columns)
sns.pairplot(scaled_df,diag_kind='kde')
```

[17]: <seaborn.axisgrid.PairGrid at 0x1d6810d5790>




```
[18]: sns.heatmap(scaled_df.corr(),annot=True,cmap='rainbow')
```


```
[18]: <Axes: >
```




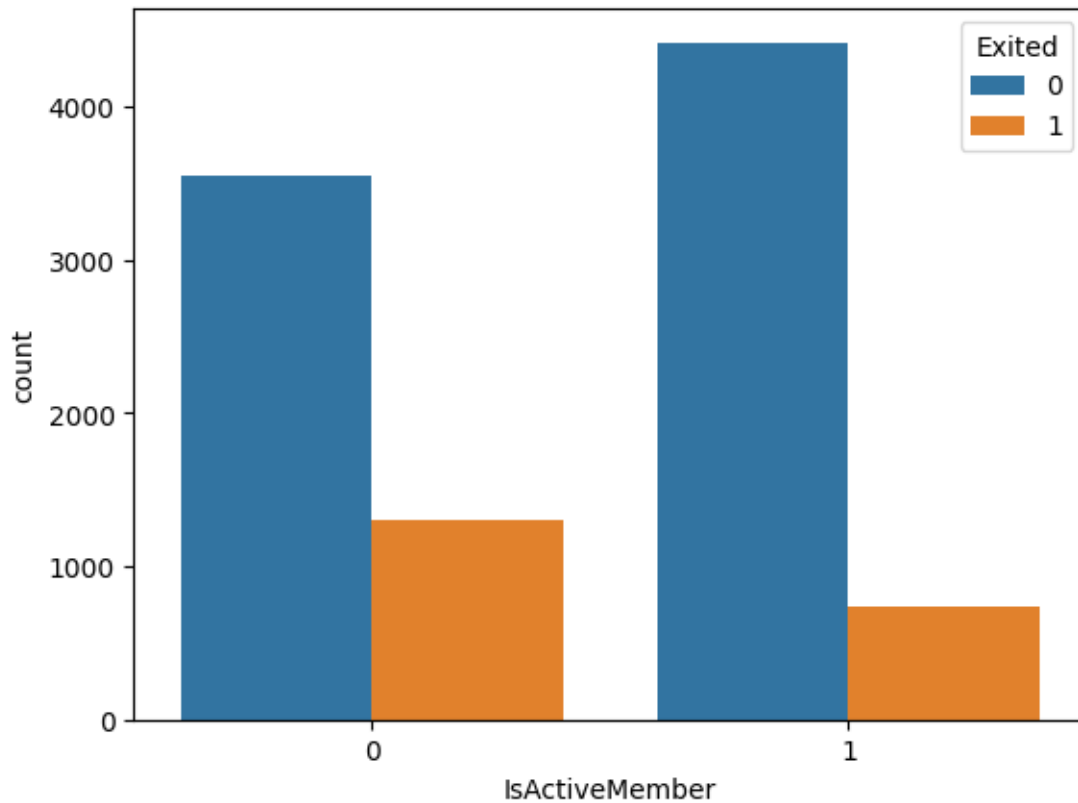
```
[22]: ## Categorical Features vs Target Variable
sns.countplot(x='Geography',data=data,hue='Exited')
plt.show()
sns.countplot(x='Gender',data=data,hue='Exited')
plt.show()
sns.countplot(x='HasCrCard',data=data,hue='Exited')
plt.show()
sns.countplot(x='IsActiveMember',data=data,hue='Exited')
plt.show()
```











```
[25]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.
↳10,random_state=5)
X_train,X_val,y_train,y_val=train_test_split(X_train,y_train,test_size=0.
↳10,random_state=5)
print("X_train size is {}".format(X_train.shape[0]))
print("X_val size is {}".format(X_val.shape[0]))
print("X_test size is {}".format(X_test.shape[0]))
```

```
X_train size is 8100
X_val size is 900
X_test size is 1000
```

```
[26]: from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
num_cols=['CreditScore','Age','Tenure','Balance','NumOfProducts','EstimatedSalary']
num_subset=scaler.fit_transform(X_train[num_cols])
X_train_num_df=pd.DataFrame(num_subset,columns=num_cols)
X_train_num_df['Geography']=list(X_train['Geography'])
X_train_num_df['Gender']=list(X_train['Gender'])
X_train_num_df['HasCrCard']=list(X_train['HasCrCard'])
```

```
X_train_num_df['IsActiveMember']=list(X_train['IsActiveMember'])
X_train_num_df.head()
```

```
[26]:
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	EstimatedSalary	\
0	-1.178587	-1.041960	-1.732257	0.198686	0.820905	1.560315	
1	-0.380169	-1.326982	1.730718	-0.022020	-0.907991	-0.713592	
2	-0.349062	1.808258	-0.693364	0.681178	0.820905	-1.126515	
3	0.625629	2.378302	-0.347067	-1.229191	0.820905	-1.682740	
4	-0.203895	-1.136967	1.730718	0.924256	-0.907991	1.332535	

	Geography	Gender	HasCrCard	IsActiveMember
0	France	Male	1	1
1	Spain	Female	1	0
2	Germany	Female	1	0
3	France	Male	1	1
4	Spain	Male	1	1

```
[27]: ## Standardise the Validation data
num_subset=scaler.fit_transform(X_val[num_cols])
X_val_num_df=pd.DataFrame(num_subset,columns=num_cols)
X_val_num_df['Geography']=list(X_val['Geography'])
X_val_num_df['Gender']=list(X_val['Gender'])
X_val_num_df['HasCrCard']=list(X_val['HasCrCard'])
X_val_num_df['IsActiveMember']=list(X_val['IsActiveMember'])
```

```
[28]: ## Standardise the Test data
num_subset=scaler.fit_transform(X_test[num_cols])
X_test_num_df=pd.DataFrame(num_subset,columns=num_cols)
X_test_num_df['Geography']=list(X_test['Geography'])
X_test_num_df['Gender']=list(X_test['Gender'])
X_test_num_df['HasCrCard']=list(X_test['HasCrCard'])
X_test_num_df['IsActiveMember']=list(X_test['IsActiveMember'])
```

```
[29]: ## Convert the categorical features to numerical
X_train_num_df=pd.get_dummies(X_train_num_df,columns=['Geography','Gender'])
X_test_num_df=pd.get_dummies(X_test_num_df,columns=['Geography','Gender'])
X_val_num_df=pd.get_dummies(X_val_num_df,columns=['Geography','Gender'])
X_train_num_df.head()
```

```
[29]:
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	EstimatedSalary	\
0	-1.178587	-1.041960	-1.732257	0.198686	0.820905	1.560315	
1	-0.380169	-1.326982	1.730718	-0.022020	-0.907991	-0.713592	
2	-0.349062	1.808258	-0.693364	0.681178	0.820905	-1.126515	
3	0.625629	2.378302	-0.347067	-1.229191	0.820905	-1.682740	
4	-0.203895	-1.136967	1.730718	0.924256	-0.907991	1.332535	

	HasCrCard	IsActiveMember	Geography_France	Geography_Germany	\
--	-----------	----------------	------------------	-------------------	---

0	1	1	True	False
1	1	0	False	False
2	1	0	False	True
3	1	1	True	False
4	1	1	False	False

	Geography_Spain	Gender_Female	Gender_Male
0	False	False	True
1	True	True	False
2	False	True	False
3	False	False	True
4	True	False	True

```
[30]: from tensorflow.keras import Sequential
      from tensorflow.keras.layers import Dense
```

```
model=Sequential()
model.add(Dense(7,activation='relu'))
model.add(Dense(10,activation='relu'))
model.add(Dense(1,activation='sigmoid'))
```

```
[31]: import tensorflow as tf
      optimizer=tf.keras.optimizers.Adam(0.01)
      model.
      ↪compile(loss='binary_crossentropy',optimizer=optimizer,metrics=['accuracy'])
```

```
[32]: model.fit(X_train_num_df,y_train,epochs=100,batch_size=10,verbose=1)
```

```
Epoch 1/100
810/810          2s 981us/step -
accuracy: 0.8120 - loss: 0.4428
Epoch 2/100
810/810          1s 1ms/step -
accuracy: 0.8546 - loss: 0.3542
Epoch 3/100
810/810          1s 797us/step -
accuracy: 0.8493 - loss: 0.3578
Epoch 4/100
810/810          1s 879us/step -
accuracy: 0.8586 - loss: 0.3458
Epoch 5/100
810/810          1s 1ms/step -
accuracy: 0.8568 - loss: 0.3542
Epoch 6/100
810/810          1s 909us/step -
accuracy: 0.8638 - loss: 0.3400
Epoch 7/100
```

810/810                    1s 1ms/step -  
 accuracy: 0.8610 - loss: 0.3436  
 Epoch 8/100  
 810/810                    1s 976us/step -  
 accuracy: 0.8689 - loss: 0.3364  
 Epoch 9/100  
 810/810                    1s 768us/step -  
 accuracy: 0.8661 - loss: 0.3352  
 Epoch 10/100  
 810/810                    1s 772us/step -  
 accuracy: 0.8691 - loss: 0.3245  
 Epoch 11/100  
 810/810                    1s 1ms/step -  
 accuracy: 0.8708 - loss: 0.3289  
 Epoch 12/100  
 810/810                    1s 865us/step -  
 accuracy: 0.8663 - loss: 0.3312  
 Epoch 13/100  
 810/810                    1s 794us/step -  
 accuracy: 0.8698 - loss: 0.3249  
 Epoch 14/100  
 810/810                    1s 725us/step -  
 accuracy: 0.8612 - loss: 0.3397  
 Epoch 15/100  
 810/810                    1s 716us/step -  
 accuracy: 0.8636 - loss: 0.3346  
 Epoch 16/100  
 810/810                    1s 957us/step -  
 accuracy: 0.8636 - loss: 0.3314  
 Epoch 17/100  
 810/810                    1s 796us/step -  
 accuracy: 0.8633 - loss: 0.3295  
 Epoch 18/100  
 810/810                    1s 994us/step -  
 accuracy: 0.8613 - loss: 0.3302  
 Epoch 19/100  
 810/810                    1s 884us/step -  
 accuracy: 0.8625 - loss: 0.3317  
 Epoch 20/100  
 810/810                    1s 979us/step -  
 accuracy: 0.8763 - loss: 0.3156  
 Epoch 21/100  
 810/810                    1s 851us/step -  
 accuracy: 0.8660 - loss: 0.3360  
 Epoch 22/100  
 810/810                    1s 740us/step -  
 accuracy: 0.8607 - loss: 0.3388  
 Epoch 23/100

810/810                    1s 764us/step -  
 accuracy: 0.8632 - loss: 0.3390  
 Epoch 24/100  
 810/810                    1s 768us/step -  
 accuracy: 0.8628 - loss: 0.3294  
 Epoch 25/100  
 810/810                    1s 717us/step -  
 accuracy: 0.8662 - loss: 0.3315  
 Epoch 26/100  
 810/810                    1s 734us/step -  
 accuracy: 0.8685 - loss: 0.3291  
 Epoch 27/100  
 810/810                    1s 754us/step -  
 accuracy: 0.8678 - loss: 0.3241  
 Epoch 28/100  
 810/810                    1s 710us/step -  
 accuracy: 0.8679 - loss: 0.3330  
 Epoch 29/100  
 810/810                    1s 1ms/step -  
 accuracy: 0.8713 - loss: 0.3242  
 Epoch 30/100  
 810/810                    1s 812us/step -  
 accuracy: 0.8719 - loss: 0.3200  
 Epoch 31/100  
 810/810                    1s 738us/step -  
 accuracy: 0.8629 - loss: 0.3244  
 Epoch 32/100  
 810/810                    1s 754us/step -  
 accuracy: 0.8655 - loss: 0.3283  
 Epoch 33/100  
 810/810                    1s 850us/step -  
 accuracy: 0.8708 - loss: 0.3211  
 Epoch 34/100  
 810/810                    1s 869us/step -  
 accuracy: 0.8720 - loss: 0.3163  
 Epoch 35/100  
 810/810                    1s 977us/step -  
 accuracy: 0.8631 - loss: 0.3313  
 Epoch 36/100  
 810/810                    1s 999us/step -  
 accuracy: 0.8699 - loss: 0.3218  
 Epoch 37/100  
 810/810                    1s 931us/step -  
 accuracy: 0.8707 - loss: 0.3212  
 Epoch 38/100  
 810/810                    1s 909us/step -  
 accuracy: 0.8721 - loss: 0.3168  
 Epoch 39/100



810/810                    1s 910us/step -  
 accuracy: 0.8721 - loss: 0.3207  
 Epoch 40/100  
 810/810                    1s 792us/step -  
 accuracy: 0.8693 - loss: 0.3210  
 Epoch 41/100  
 810/810                    1s 1ms/step -  
 accuracy: 0.8569 - loss: 0.3405  
 Epoch 42/100  
 810/810                    1s 762us/step -  
 accuracy: 0.8746 - loss: 0.3176  
 Epoch 43/100  
 810/810                    1s 680us/step -  
 accuracy: 0.8696 - loss: 0.3259  
 Epoch 44/100  
 810/810                    1s 711us/step -  
 accuracy: 0.8611 - loss: 0.3351  
 Epoch 45/100  
 810/810                    1s 867us/step -  
 accuracy: 0.8675 - loss: 0.3214  
 Epoch 46/100  
 810/810                    1s 746us/step -  
 accuracy: 0.8671 - loss: 0.3205  
 Epoch 47/100  
 810/810                    1s 763us/step -  
 accuracy: 0.8725 - loss: 0.3174  
 Epoch 48/100  
 810/810                    1s 695us/step -  
 accuracy: 0.8725 - loss: 0.3233  
 Epoch 49/100  
 810/810                    1s 643us/step -  
 accuracy: 0.8685 - loss: 0.3310  
 Epoch 50/100  
 810/810                    1s 655us/step -  
 accuracy: 0.8747 - loss: 0.3152  
 Epoch 51/100  
 810/810                    1s 849us/step -  
 accuracy: 0.8647 - loss: 0.3288  
 Epoch 52/100  
 810/810                    1s 1ms/step -  
 accuracy: 0.8663 - loss: 0.3190  
 Epoch 53/100  
 810/810                    1s 926us/step -  
 accuracy: 0.8729 - loss: 0.3101  
 Epoch 54/100  
 810/810                    1s 722us/step -  
 accuracy: 0.8688 - loss: 0.3257  
 Epoch 55/100

810/810                    1s 812us/step -  
 accuracy: 0.8653 - loss: 0.3206  
 Epoch 56/100  
 810/810                    1s 734us/step -  
 accuracy: 0.8717 - loss: 0.3258  
 Epoch 57/100  
 810/810                    1s 699us/step -  
 accuracy: 0.8616 - loss: 0.3270  
 Epoch 58/100  
 810/810                    1s 817us/step -  
 accuracy: 0.8643 - loss: 0.3231  
 Epoch 59/100  
 810/810                    1s 1ms/step -  
 accuracy: 0.8761 - loss: 0.3093  
 Epoch 60/100  
 810/810                    1s 1ms/step -  
 accuracy: 0.8668 - loss: 0.3252  
 Epoch 61/100  
 810/810                    1s 1ms/step -  
 accuracy: 0.8637 - loss: 0.3276  
 Epoch 62/100  
 810/810                    1s 1ms/step -  
 accuracy: 0.8688 - loss: 0.3241  
 Epoch 63/100  
 810/810                    1s 1ms/step -  
 accuracy: 0.8644 - loss: 0.3307  
 Epoch 64/100  
 810/810                    1s 1ms/step -  
 accuracy: 0.8658 - loss: 0.3240  
 Epoch 65/100  
 810/810                    1s 992us/step -  
 accuracy: 0.8745 - loss: 0.3198  
 Epoch 66/100  
 810/810                    1s 961us/step -  
 accuracy: 0.8586 - loss: 0.3365  
 Epoch 67/100  
 810/810                    1s 944us/step -  
 accuracy: 0.8682 - loss: 0.3162  
 Epoch 68/100  
 810/810                    1s 808us/step -  
 accuracy: 0.8675 - loss: 0.3207  
 Epoch 69/100  
 810/810                    1s 904us/step -  
 accuracy: 0.8621 - loss: 0.3342  
 Epoch 70/100  
 810/810                    1s 930us/step -  
 accuracy: 0.8586 - loss: 0.3359  
 Epoch 71/100

810/810                    1s 1ms/step -  
 accuracy: 0.8706 - loss: 0.3193  
 Epoch 72/100  
 810/810                    1s 1ms/step -  
 accuracy: 0.8635 - loss: 0.3239  
 Epoch 73/100  
 810/810                    1s 885us/step -  
 accuracy: 0.8756 - loss: 0.3081  
 Epoch 74/100  
 810/810                    1s 1ms/step -  
 accuracy: 0.8602 - loss: 0.3381  
 Epoch 75/100  
 810/810                    1s 1ms/step -  
 accuracy: 0.8643 - loss: 0.3317  
 Epoch 76/100  
 810/810                    1s 974us/step -  
 accuracy: 0.8730 - loss: 0.3187  
 Epoch 77/100  
 810/810                    1s 987us/step -  
 accuracy: 0.8662 - loss: 0.3264  
 Epoch 78/100  
 810/810                    1s 978us/step -  
 accuracy: 0.8731 - loss: 0.3162  
 Epoch 79/100  
 810/810                    1s 1ms/step -  
 accuracy: 0.8751 - loss: 0.3123  
 Epoch 80/100  
 810/810                    1s 1ms/step -  
 accuracy: 0.8678 - loss: 0.3264  
 Epoch 81/100  
 810/810                    1s 939us/step -  
 accuracy: 0.8638 - loss: 0.3301  
 Epoch 82/100  
 810/810                    1s 1ms/step -  
 accuracy: 0.8692 - loss: 0.3203  
 Epoch 83/100  
 810/810                    1s 1ms/step -  
 accuracy: 0.8668 - loss: 0.3263  
 Epoch 84/100  
 810/810                    1s 1ms/step -  
 accuracy: 0.8653 - loss: 0.3251  
 Epoch 85/100  
 810/810                    1s 1ms/step -  
 accuracy: 0.8711 - loss: 0.3169  
 Epoch 86/100  
 810/810                    1s 1ms/step -  
 accuracy: 0.8674 - loss: 0.3206  
 Epoch 87/100

```

810/810          1s 971us/step -
accuracy: 0.8783 - loss: 0.3129
Epoch 88/100
810/810          1s 1ms/step -
accuracy: 0.8683 - loss: 0.3225
Epoch 89/100
810/810          1s 1ms/step -
accuracy: 0.8721 - loss: 0.3165
Epoch 90/100
810/810          1s 1ms/step -
accuracy: 0.8647 - loss: 0.3268
Epoch 91/100
810/810          1s 1ms/step -
accuracy: 0.8768 - loss: 0.3166
Epoch 92/100
810/810          1s 1ms/step -
accuracy: 0.8683 - loss: 0.3260
Epoch 93/100
810/810          1s 1ms/step -
accuracy: 0.8661 - loss: 0.3276
Epoch 94/100
810/810          1s 1ms/step -
accuracy: 0.8708 - loss: 0.3194
Epoch 95/100
810/810          1s 1ms/step -
accuracy: 0.8576 - loss: 0.3367
Epoch 96/100
810/810          1s 1ms/step -
accuracy: 0.8719 - loss: 0.3115
Epoch 97/100
810/810          1s 1ms/step -
accuracy: 0.8663 - loss: 0.3268
Epoch 98/100
810/810          1s 1ms/step -
accuracy: 0.8712 - loss: 0.3151
Epoch 99/100
810/810          1s 1ms/step -
accuracy: 0.8613 - loss: 0.3282
Epoch 100/100
810/810          1s 1ms/step -
accuracy: 0.8675 - loss: 0.3260

```

[32]: <keras.src.callbacks.history.History at 0x1d69617b200>

```

[34]: from sklearn.metrics import confusion_matrix
y_pred_val = model.predict(X_val)  # Replace model with your trained model
    ↪variable

```

```

y_pred_val = (y_pred_val > 0.5).astype(int) # Convert probabilities to binary
↳ predictions

# Now you can compute the confusion matrix
cm_val = confusion_matrix(y_val, y_pred_val)
print("Confusion Matrix:")
print(cm_val)

# Optional: Print the classification report for more detailed evaluation
print("\nClassification Report:")
print(classification_report(y_val, y_pred_val))

```

```

-----
ValueError                                Traceback (most recent call last)
Cell In[34], line 2
      1 from sklearn.metrics import confusion_matrix
----> 2 y_pred_val = model.predict(X_val) # Replace model with your trained
↳ model variable
      3 y_pred_val = (y_pred_val > 0.5).astype(int) # Convert probabilities to
↳ binary predictions
      5 # Now you can compute the confusion matrix

File
↳ ~\AppData\Roaming\Python\Python312\site-packages\keras\src\utils\traceback_utils.py:122, in filter_traceback.<locals>.error_handler(*args, **kwargs)
    119     filtered_tb = _process_traceback_frames(e.__traceback__)
    120     # To get the full stack trace, call:
    121     # `keras.config.disable_traceback_filtering()`
--> 122     raise e.with_traceback(filtered_tb) from None
    123 finally:
    124     del filtered_tb

File ~\AppData\Roaming\Python\Python312\site-packages\optree\ops.py:747, in
↳ tree_map(func, tree, is_leaf, none_is_leaf, namespace, *rests)
    745 leaves, treespec = _C.flatten(tree, is_leaf, none_is_leaf, namespace)
    746 flat_args = [leaves] + [treespec.flatten_up_to(r) for r in rests]
--> 747 return treespec.unflatten(map(func, *flat_args))

File ~\AppData\Roaming\Python\Python312\site-packages\pandas\core\generic.py:
↳ 6640, in NDFrame.astype(self, dtype, copy, errors)
    6634     results = [
    6635         ser.astype(dtype, copy=copy, errors=errors) for _, ser in self.
↳ items()
    6636     ]
    6638 else:
    6639     # else, only a single dtype is given
-> 6640     new_data = self._mgr.astype(dtype=dtype, copy=copy, errors=errors)

```

```

6641     res = self._constructor_from_mgr(new_data, axes=new_data.axes)
6642     return res.__finalize__(self, method="astype")

```

File

```

→ ~\AppData\Roaming\Python\Python312\site-packages\pandas\core\internals\managers.py
→ py:430, in BaseBlockManager.astype(self, dtype, copy, errors)
    427 elif using_copy_on_write():
    428     copy = False
--> 430 return self.apply(
    431     "astype",
    432     dtype=dtype,
    433     copy=copy,
    434     errors=errors,
    435     using_cow=using_copy_on_write(),
    436 )

```

File

```

→ ~\AppData\Roaming\Python\Python312\site-packages\pandas\core\internals\managers.py
→ py:363, in BaseBlockManager.apply(self, f, align_keys, **kwargs)
    361     applied = b.apply(f, **kwargs)
    362     else:
--> 363     applied = getattr(b, f)(**kwargs)
    364     result_blocks = extend_blocks(applied, result_blocks)
    366 out = type(self).from_blocks(result_blocks, self.axes)

```

File

```

→ ~\AppData\Roaming\Python\Python312\site-packages\pandas\core\internals\blocks.py
→ py:758, in Block.astype(self, dtype, copy, errors, using_cow, squeeze)
    755     raise ValueError("Can not squeeze with more than one column.")
    756     values = values[0, :] # type: ignore[call-overload]
--> 758 new_values = astype_array_safe(values, dtype, copy=copy, errors=errors)
    760 new_values = maybe_coerce_values(new_values)
    762 refs = None

```

File ~\AppData\Roaming\Python\Python312\site-packages\pandas\core\dtypes\astype

```

→ py:237, in astype_array_safe(values, dtype, copy, errors)
    234     dtype = dtype.numpy_dtype
    236 try:
--> 237     new_values = astype_array(values, dtype, copy=copy)
    238 except (ValueError, TypeError):
    239     # e.g. _astype_nansafe can fail on object-dtype of strings
    240     # trying to convert to float
    241     if errors == "ignore":

```

File ~\AppData\Roaming\Python\Python312\site-packages\pandas\core\dtypes\astype

```

→ py:182, in astype_array(values, dtype, copy)
    179     values = values.astype(dtype, copy=copy)
    181 else:

```

```

--> 182     values = _astype_nansafe(values, dtype, copy=copy)
      184 # in pandas we don't store numpy str dtypes, so convert to object
      185 if isinstance(dtype, np.dtype) and issubclass(values.dtype.type, str):

File ~\AppData\Roaming\Python\Python312\site-packages\pandas\core\dtypes\astype
py:133, in _astype_nansafe(arr, dtype, copy, skipna)
      129     raise ValueError(msg)
      131 if copy or arr.dtype == object or dtype == object:
      132     # Explicit copy, or required since NumPy can't view from / to objec
--> 133     return arr.astype(dtype, copy=True)
      135 return arr.astype(dtype, copy=copy)

```

**ValueError:** could not convert string to float: 'Kerr'

```

[ ]: Accuracy=782/900
print("Accuracy of the Model on the Validation Data set is 86.89%")

```

```

[ ]: loss1,accuracy1=model.evaluate(X_train_num_df,y_train,verbose=False)
loss2,accuracy2=model.evaluate(X_val_num_df,y_val,verbose=False)
print("Train Loss {}".format(loss1))
print("Train Accuracy {}".format(accuracy1))
print("Val Loss {}".format(loss2))
print("Val Accuracy {}".format(accuracy2))

```

```

[ ]:

```

# exp-ml-4

October 14, 2024

```
[1]: import numpy as np
import matplotlib.pyplot as plt
```

```
[2]: def gradient_descent(learning_rate, max_iterations, initial_x):
    x = initial_x
    x_history = [] # Create a list to store the history of x values
    for _ in range(max_iterations):
        gradient = 2 * (x + 3) # Compute the gradient of the function
        x = x - learning_rate * gradient # Update x using the gradient and
        ↪ learning rate
        x_history.append(x) # Append the current x to the history list

    return x, x_history
```

```
[5]: # Parameters for Gradient Descent
learning_rate = 0.1
max_iterations = 1000
initial_x = 2

# Run Gradient Descent to find the local minimum
local_minimum, x_history = gradient_descent(learning_rate, max_iterations,
    ↪ initial_x)

print(f"Local Minimum at x = {local_minimum}")


# Plot the graph to visualize the convergence
x_values = np.linspace(-10, 10, 400) # Generate x values for the graph
y_values = (x_values + 3)**2 # Calculate corresponding y values

plt.plot(x_values, y_values, label='y = (x + 3)^2', color='blue')
plt.scatter(x_history, [(x + 3)**2 for x in x_history], label='Gradient Descent',
    ↪ Path', color='red', marker='x')
plt.xlabel('x')
plt.ylabel('y')
plt.legend()
plt.title('Gradient Descent Convergence')
plt.grid(True)
```



```
plt.show()
```

Local Minimum at  $x = -2.9999999999999999$



```
[ ]:
```

# exp-ml-5

October 14, 2024

```
[1]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, \
recall_score, f1_score
```

```
[2]: data=pd.read_csv("diabetes.csv")
data
```

```
[2]:      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
0                6      148                72              35         0  33.6
1                1       85                66              29         0  26.6
2                8      183                64               0         0  23.3
3                1       89                66              23        94  28.1
4                0      137                40              35       168  43.1
..          ...    ...          ...          ...    ...    ...
763             10      101                76              48       180  32.9
764              2      122                70              27         0  36.8
765              5      121                72              23       112  26.2
766              1      126                60               0         0  30.1
767              1       93                70              31         0  30.4
```

```
      Pedigree  Age  Outcome
0      0.627   50         1
1      0.351   31         0
2      0.672   32         1
3      0.167   21         0
4      2.288   33         1
..          ...  ...    ...
763     0.171   63         0
764     0.340   27         0
765     0.245   30         0
766     0.349   47         1
767     0.315   23         0
```

[768 rows x 9 columns]

```
[3]: X = data.drop("Outcome", axis=1) # Features
      y = data["Outcome"] # Target variable
```

```
[4]: X
```

```
[4]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
..	...	...	...	...	...	...	
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

	Pedigree	Age
0	0.627	50
1	0.351	31
2	0.672	32
3	0.167	21
4	2.288	33
..	...	...
763	0.171	63
764	0.340	27
765	0.245	30
766	0.349	47
767	0.315	23

[768 rows x 8 columns]

```
[5]: # 2. Split the dataset into training and test sets
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪ random_state=42)
```

```
[6]: # 3. Normalize the data
      scaler = StandardScaler()
      X_train = scaler.fit_transform(X_train)
      X_test = scaler.transform(X_test)
```

```
[7]: X_train
```

```
[7]: array([[ -0.52639686, -1.15139792, -3.75268255, ..., -4.13525578,
          -0.49073479, -1.03594038],
          [ 1.58804586, -0.27664283,  0.68034485, ..., -0.48916881,
           2.41502991,  1.48710085],
          [-0.82846011,  0.56687102, -1.2658623 , ..., -0.42452187,
           0.54916055, -0.94893896],
          ...,
          [ 1.8901091 , -0.62029661,  0.89659009, ...,  1.76054443,
           1.981245  ,  0.44308379],
          [-1.13052335,  0.62935353, -3.75268255, ...,  1.34680407,
          -0.78487662, -0.33992901],
          [-1.13052335,  0.12949347,  1.43720319, ..., -1.22614383,
          -0.61552223, -1.03594038]])
```

```
[8]: # 4. Implement K-Nearest Neighbors (KNN)
k = 3 # Choose the number of neighbors (k) based on your needs
knn = KNeighborsClassifier(n_neighbors=k)
knn.fit(X_train, y_train)
```

```
[8]: KNeighborsClassifier(n_neighbors=3)
```

```
[11]: # 5. Predict and Evaluate
y_pred = knn.predict(X_test)
y_pred
```

```
[11]: array([0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0,
           0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0,
           0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 1,
           0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0,
           0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1,
           0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1,
           0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0],
          dtype=int64)
```

```
[12]: from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, \
      ↪ recall_score

# Assuming y_test and y_pred are already defined
conf_matrix = confusion_matrix(y_test, y_pred)

# Calculate accuracy, error rate, precision, and recall
accuracy = accuracy_score(y_test, y_pred)
error_rate = 1 - accuracy
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)

# Print results
```

```
print("Confusion Matrix:")
print(conf_matrix)
print("Accuracy:", accuracy)
print("Error Rate:", error_rate)
print("Precision:", precision)
print("Recall:", recall)
```

Confusion Matrix:

```
[[81 18]
 [27 28]]
```

Accuracy: 0.7077922077922078

Error Rate: 0.29220779220779225

Precision: 0.6086956521739131

Recall: 0.509090909090909